

Challenges for Gaussian processes

Mark van der Wilk

Department of Computing
Imperial College London

 @markvanderwilk
m.vdwilk@imperial.ac.uk

February 3, 2023

Methods for Regression in ML

Most tasks in ML are regression! Deep learning is 99% regression!

Remember where GPs excel:

- ▶ are low dimensional (e.g. tens of dimensions rather than 100s),
- ▶ have little data (or data is expensive to obtain),
- ▶ are noisy (random fluctuations that obscure the signal),
- ▶ require uncertainty estimates.

Today: Where GPs struggle, and why.

Learning objectives

Know how to perform the computations necessary for GPs

- ▶ Computing the marginal likelihood
- ▶ Computational complexity
- ▶ Low-rank approximations

Understand the computational and modelling limitations of GPs

- ▶ Limitations of stationary kernels
- ▶ Limitations of local kernels in high-dimensions

Training a GP: Computations

To train, we need the marginal likelihood and its gradient:

$$\log p(\mathbf{y} | \theta) = \text{const} - \frac{1}{2} \log |\mathbf{K}_\theta + \sigma^2 \mathbf{I}| - \frac{1}{2} \mathbf{y}^\top (\mathbf{K}_\theta + \sigma^2 \mathbf{I})^{-1} \mathbf{y} \quad (1)$$

$$\frac{\partial}{\partial \theta} \log p(\mathbf{y} | \theta) = \frac{1}{2} \text{Tr} \left[\left(\mathbf{K}_\theta^{-1} \mathbf{y} \mathbf{y}^\top \mathbf{K}_\theta - \mathbf{K}_\theta^{-1} \right) \frac{\partial \mathbf{K}_\theta}{\partial \theta} \right] \quad (2)$$

$(\dots)^{-1}$ and $|\dots|$ are calculated from a **matrix decomposition**.

- ▶ Decompositions are expensive
- ▶ But make follow-on operations cheap
- ▶ The correct decomposition helps numerical stability¹
- ▶ Directly computing the inverse (`np.linalg.inv()`) is a bad thing to do!

¹<https://nhigham.com/2020/08/04/what-is-numerical-stability/>

Eigenvalue decomposition

We can compute the terms using the **eigenvalue decomposition**:

$$\mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^\top = \mathbf{K}_\theta + \sigma^2\mathbf{I} \quad (3)$$

$$\log|\mathbf{K}_\theta + \sigma^2\mathbf{I}| = \log|\mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^\top| = \log|\mathbf{Q}| + \log|\mathbf{\Lambda}| = \sum_{n=1}^N \log \lambda_n(\mathbf{K}_\theta + \sigma^2\mathbf{I}) \quad (4)$$

$$\mathbf{y}^\top (\mathbf{K}_\theta + \sigma^2\mathbf{I})^{-1} \mathbf{y} = \mathbf{y}^\top \mathbf{Q}^\top \mathbf{\Lambda}^{-1} \mathbf{Q} \mathbf{y} \quad (5)$$

Eigenvalue decomposition is mostly applied for **theoretical analysis**.

Cholesky decomposition

Or alternatively the **Cholesky decomposition**:

$$\mathbf{L}\mathbf{L}^\top = \mathbf{K}_\theta + \sigma^2\mathbf{I}, \quad \text{where } \mathbf{L} \text{ is lower triangular.} \quad (6)$$

$$\log|\mathbf{K}_\theta + \sigma^2\mathbf{I}| = \log|\mathbf{L}\mathbf{L}^\top| = 2 \sum_{n=1}^N \log[\mathbf{L}]_{nn} \quad (7)$$

$$\mathbf{y}^\top (\mathbf{K}_\theta + \sigma^2\mathbf{I})^{-1} \mathbf{y} = \mathbf{y}^\top \mathbf{L}^\top{}^{-1} \mathbf{L}^{-1} \mathbf{y} \quad (8)$$

Cholesky decomposition is used for **practical implementation**.
(For the coursework, all implementations are fine.)

Computational complexity

1. Computing kernel matrix $O(N^2)$ time and $O(N^2)$ space
2. Eigendecomposition and Cholesky are both often quoted to be $O(N^3)$ time.²
3. Logdet or inverse given the decomposition are fast
 - ▶ Logdet for both are $O(N)$
 - ▶ Inverse is $O(N^2)$

Can we take advantage of structure in the kernel matrix to do better?

²Algorithms do exist with better asymptotic complexity, but they are definitely slower than $O(N^2)$.

Low-rank kernels

For kernel matrices that are low-rank (have some zero eigenvalues),
i.e. with $M \ll N$

$$\mathbf{K} = \mathbf{P}\mathbf{P}^\top, \quad \text{where } \mathbf{K} \in \mathbb{R}^{N \times N}, \mathbf{P} \in \mathbb{R}^{N \times M} \quad (9)$$

we can compute the marginal likelihood more cheaply:³

$$\log|\mathbf{K} + \sigma^2\mathbf{I}_N| = \log|\sigma^2\mathbf{I}_N| + \log|\mathbf{I}_M + \sigma^{-2}\mathbf{P}^\top\mathbf{P}| \quad (10)$$

$$\mathbf{y}^\top(\mathbf{K} + \sigma^2\mathbf{I}_N)^{-1}\mathbf{y} = \mathbf{y}^\top\left(\sigma^{-2}\mathbf{I}_N - \sigma^{-2}\mathbf{P}(\sigma^2\mathbf{I}_M + \mathbf{P}^\top\mathbf{P})^{-1}\mathbf{P}^\top\right)\mathbf{y} \quad (11)$$

Can be computed in $O(NM^2 + M^3)$.

³Can be proved by applying the Woodbury Matrix Identity, and the similar Matrix Determinant Lemma.

Decaying eigenvalues

For example, the squared exponential kernel

$$k(x, x') = \sigma_f^2 \exp\left(\frac{|x - x'|^2}{2\ell^2}\right), \quad \theta = \{\sigma_f, \ell\} \quad (12)$$

with 1D inputs i.i.d. from⁴ $\mathcal{N}(0, \sigma_D^2)$, we have that the eigenvalues of \mathbf{K} decay exponentially

$$\lambda_m = \sqrt{\frac{2a}{A}} B^m \quad \text{with } B < 1 \quad (13)$$

$$a^{-1} = 4\sigma_D^2, \quad b^{-1} = 2\ell^2, \quad c = \sqrt{a^2 + 2ab}, \quad A = a + b + c, \quad B = b/A$$

as the number of data $N \rightarrow \infty$.

⁴Conditions can be much weaker, but this is easier.

Low-rank approximations

So we know that for large N our \mathbf{K} will be **approximately low-rank**. Can we approximate it with a low-rank matrix?

$$\mathbf{K} = \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^\top = [\mathbf{Q}_1 \quad \mathbf{Q}_2] \begin{bmatrix} \mathbf{\Lambda}_1 & \\ & \mathbf{\Lambda}_2 \end{bmatrix} \begin{bmatrix} \mathbf{Q}_1^\top \\ \mathbf{Q}_2^\top \end{bmatrix} = \mathbf{Q}_1\mathbf{\Lambda}_1\mathbf{Q}_1^\top + \mathbf{Q}_2\mathbf{\Lambda}_2\mathbf{Q}_2^\top \quad (14)$$

$$\hat{\mathbf{K}} = \mathbf{Q}_1\mathbf{\Lambda}_1\mathbf{Q}_1^\top + \cancel{\mathbf{Q}_2\mathbf{\Lambda}_2\mathbf{Q}_2^\top} \stackrel{\approx 0}{=} \mathbf{P}\mathbf{P}^\top, \quad \text{with } \mathbf{P} = \mathbf{Q}_1\mathbf{\Lambda}_1^{\frac{1}{2}}. \quad (15)$$

This gives the approximation with the minimal Frobenius norm:

$$\|\mathbf{K} - \hat{\mathbf{K}}\|_F = \sqrt{\sum_{i=M+1}^N \lambda_i} \quad (16)$$

Low-rank approximation: Data fit

Starting with the data-fit term:⁵

$$\begin{aligned}\mathbf{y}^\top (\mathbf{K} + \sigma^2 \mathbf{I})^{-1} \mathbf{y} &= \mathbf{y}^\top (\hat{\mathbf{K}} + \sigma^2 \mathbf{I} + \mathbf{K} - \hat{\mathbf{K}})^{-1} \mathbf{y} \\ &= \mathbf{y}^\top (\hat{\mathbf{K}} + \sigma^2 \mathbf{I})^{-1} \mathbf{y} - \underbrace{\mathbf{y}^\top (\hat{\mathbf{K}} + \sigma^2 \mathbf{I})^{-1} (\mathbf{K} - \hat{\mathbf{K}}) (\mathbf{K} + \sigma^2 \mathbf{I})^{-1} \mathbf{y}}_{\rightarrow 0 \text{ as } \|\mathbf{K} - \hat{\mathbf{K}}\|_F \rightarrow 0} \quad (17)\end{aligned}$$

You can place bounds on the error using knowledge of λ_M (the final included eigenvalue).

⁵Applying the identity $(A + B)^{-1} = A^{-1} - A^{-1}B(A + B)^{-1}$.

Low-rank approximation: Complexity penalty

And finishing with the complexity penalty:

$$\begin{aligned}\log|\mathbf{K} + \sigma^2\mathbf{I}| &= \sum_{n=1}^N \log(\lambda_n(\mathbf{K}) + \sigma^2) \\ &= \sum_{m=1}^M \log(\lambda_m(\mathbf{K}) + \sigma^2) + \sum_{r=M+1}^N \underbrace{\log(\lambda_r(\mathbf{K}) + \sigma^2)}_{\approx \sigma^2} \quad (18)\end{aligned}$$

$$\approx \sum_{m=1}^M \log(\lambda_m(\mathbf{K}) + \sigma^2) + \sum_{r=M+1}^N \log(\sigma^2) \quad (19)$$

$$= \log|\hat{\mathbf{K}} + \sigma^2\mathbf{I}| \quad (20)$$

Low-rank approximations overview

- ▶ Learning with exact matrix decompositions is expensive — $O(N^3)$ time, $O(N^2)$ memory
- ▶ Low-rank kernels improve things when $N \gg M$ — $O(NM^2)$ time, $O(NM)$ memory
- ▶ Low-rank approximations do exist!
- ▶ We used an impractical method (eigendecomposition) to find **which** low-rank columns to use, with a cost of $O(N^2M)$ time.

For certain approximations, we have 1) fast methods for finding \mathbf{P} , and 2) proofs for how large M needs to be.

E.g. Burt et al [2019] show that for regression and the Squared Exponential kernel, we can get arbitrarily good approximations in

$$O(N(\log N)^{2D}(\log \log N)^2) \quad (21)$$

$$O(N(\log N)^{2D}(\log \log N)^2) = O(N^{1+\epsilon}), \quad \forall \epsilon > 0 \quad (22)$$

Practical constraints and implementation

GPs are slow, but often not even because of asymptotic complexity!

- ▶ Memory is often the main bottleneck (storing large matrices).
This is poorly suited to modern-day hardware (i.e. GPUs)
- ▶ Decompositions require serial computations.
This is poorly suited to modern-day hardware (i.e. GPUs)
- ▶ Decompositions need high floating point precision.
This is poorly suited to modern-day hardware (i.e. GPUs)

There is hope

Approximations can make the computations look more similar to a neural network by

- ▶ Reduce the size of the matrix decomposition [Titsias, 2009]
- ▶ Improve data parallelism [Hensman et al, 2013]
- ▶ Increase reliance on fast matrix-matrix multiplications [Gardner et al, 2018; Wang et al, 2019]
- ▶ Remove matrix inverses altogether [van der Wilk et al, 2019]

This slide for reference only, not examinable.

Implementation matters!

We have good asymptotic bounds on the computational complexity!

Constants cause slowness \implies implementation matters

E.g. *Exact Gaussian Processes on a Million Data Points*, Wang et al [2019] tailors a very good approximation to GPU hardware, with impressive results.

Conclusions on Computational Issues

- ▶ Gaussian processes are computationally expensive to train
- ▶ Low-rank approximations to the kernel matrix \mathbf{K} do exist
- ▶ Low-rank approximations improve the asymptotic complexity
- ▶ Still need work to optimise algorithms and tailor to modern hardware

Modelling limitations

We know that the generalisation capability of a model depends strongly on the properties of its prior.

Do we really have good priors?

Some criticisms...

Stationary kernels

We have seen many examples of **stationary** kernels:

$$k(\mathbf{x}, \mathbf{x}') = k(\mathbf{x} - \mathbf{x}') \quad (23)$$

e.g. the squared exponential.

- ▶ Same generalisation characteristic is applied through the entire input space.
- ▶ See `stationary-kernel-failure.ipynb` in the `inference-plots` GitHub repo.

Local kernels in high dimensions

Many kernels depend on some form of the Euclidean distance:

$$k(\mathbf{x}, \mathbf{x}') = \sigma_f^2 \exp \left[-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2\ell^2} \right] \quad (24)$$

This behaves badly in high dimensions (even MNIST, $D = 28^2 = 784$):

- ▶ Small differences in each dimension add up to large distances
→ low correlation
- ▶ Function is allowed to vary independently along each input dimension
- ▶ Prior with low lengthscale is too flexible, large lengthscale is too inflexible
- ▶ No middle ground

Conclusions

Gaussian processes have limitations for both

- ▶ low-dimensional inputs (e.g. stationarity)
- ▶ high-dimensional inputs (e.g. locality)

We need to either make much better priors, or incorporate GPs into more complex models.

References I

- [1] J. Gardner, G. Pleiss, K. Q. Weinberger, D. Bindel, and A. G. Wilson. Gpytorch: Blackbox matrix-matrix gaussian process inference with gpu acceleration. In S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett, editors, *Advances in Neural Information Processing Systems 31*, pages 7576–7586. Curran Associates, Inc., 2018.