

Challenges for Gaussian processes

Mark van der Wilk

Department of Computing
Imperial College London

@markvanderwilk
m.vdwilk@imperial.ac.uk

February 4, 2022

Methods for Regression in ML

Most tasks in ML are regression! Deep learning is 99% regression!

Methods for Regression in ML

Most tasks in ML are regression! Deep learning is 99% regression!

Remember where GPs excel:

- ▶ are low dimensional (e.g. tens of dimensions rather than 100s),
- ▶ have little data (or data is expensive to obtain),
- ▶ are noisy (random fluctuations that obscure the signal),
- ▶ require uncertainty estimates.

Methods for Regression in ML

Most tasks in ML are regression! Deep learning is 99% regression!

Remember where GPs excel:

- ▶ are low dimensional (e.g. tens of dimensions rather than 100s),
- ▶ have little data (or data is expensive to obtain),
- ▶ are noisy (random fluctuations that obscure the signal),
- ▶ require uncertainty estimates.

Today: Where GPs struggle, and why.

Learning objectives

Know how to perform the computations necessary for GPs

- ▶ Computing the marginal likelihood
- ▶ Computational complexity
- ▶ Low-rank approximations

Learning objectives

Know how to perform the computations necessary for GPs

- ▶ Computing the marginal likelihood
- ▶ Computational complexity
- ▶ Low-rank approximations

Understand the computational and modelling limitations of GPs

- ▶ Limitations of stationary kernels
- ▶ Limitations of local kernels in high-dimensions

Training a GP: Computations

To train, we need the marginal likelihood and its gradient:

$$\log p(\mathbf{y} | \theta) = \text{const} - \frac{1}{2} \log |\mathbf{K}_\theta + \sigma^2 \mathbf{I}| - \frac{1}{2} \mathbf{y}^\top (\mathbf{K}_\theta + \sigma^2 \mathbf{I})^{-1} \mathbf{y} \quad (1)$$

$$\frac{\partial}{\partial \theta} \log p(\mathbf{y} | \theta) = \frac{1}{2} \text{Tr} \left[\left(\mathbf{K}_\theta^{-1} \mathbf{y} \mathbf{y}^\top \mathbf{K}_\theta - \mathbf{K}_\theta^{-1} \right) \frac{\partial \mathbf{K}_\theta}{\partial \theta} \right] \quad (2)$$

¹<https://nhigham.com/2020/08/04/what-is-numerical-stability/>

Training a GP: Computations

To train, we need the marginal likelihood and its gradient:

$$\log p(\mathbf{y} | \theta) = \text{const} - \frac{1}{2} \log |\mathbf{K}_\theta + \sigma^2 \mathbf{I}| - \frac{1}{2} \mathbf{y}^\top (\mathbf{K}_\theta + \sigma^2 \mathbf{I})^{-1} \mathbf{y} \quad (1)$$

$$\frac{\partial}{\partial \theta} \log p(\mathbf{y} | \theta) = \frac{1}{2} \text{Tr} \left[\left(\mathbf{K}_\theta^{-1} \mathbf{y} \mathbf{y}^\top \mathbf{K}_\theta - \mathbf{K}_\theta^{-1} \right) \frac{\partial \mathbf{K}_\theta}{\partial \theta} \right] \quad (2)$$

$(\dots)^{-1}$ and $|\dots|$ are calculated from a **matrix decomposition**.

- ▶ Decompositions are expensive
- ▶ But make follow-on operations cheap
- ▶ The correct decomposition helps numerical stability¹
- ▶ Directly computing the inverse (`np.linalg.inv()`) is a bad thing to do!

¹<https://nhigham.com/2020/08/04/what-is-numerical-stability/>

Eigenvalue decomposition

We can compute the terms using the **eigenvalue decomposition**:

$$\mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^\top = \mathbf{K}_\theta + \sigma^2\mathbf{I} \quad (3)$$

$$\log|\mathbf{K}_\theta + \sigma^2\mathbf{I}| = \log|\mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^\top| = \log|\mathbf{Q}|^0 + \log|\mathbf{\Lambda}| = \sum_{n=1}^N \log \lambda_i(\mathbf{K}_\theta + \sigma^2\mathbf{I}) \quad (4)$$

$$\mathbf{y}^\top (\mathbf{K}_\theta + \sigma^2\mathbf{I})^{-1} \mathbf{y} = \mathbf{y}^\top \mathbf{Q}^\top \mathbf{\Lambda}^{-1} \mathbf{Q} \mathbf{y} \quad (5)$$

Eigenvalue decomposition is mostly applied for **theoretical analysis**.

Cholesky decomposition

Or alternatively the **Cholesky decomposition**:

$$\mathbf{L}\mathbf{L}^\top = \mathbf{K}_\theta + \sigma^2\mathbf{I}, \quad \text{where } \mathbf{L} \text{ is lower triangular.} \quad (6)$$

$$\log|\mathbf{K}_\theta + \sigma^2\mathbf{I}| = \log|\mathbf{L}\mathbf{L}^\top| = 2 \sum_{n=1}^N \log[\mathbf{L}]_{nn} \quad (7)$$

$$\mathbf{y}^\top (\mathbf{K}_\theta + \sigma^2\mathbf{I})^{-1} \mathbf{y} = \mathbf{y}^\top \mathbf{L}^\top{}^{-1} \mathbf{L}^{-1} \mathbf{y} \quad (8)$$

Cholesky decomposition is used for **practical implementation**.

Cholesky decomposition

Or alternatively the **Cholesky decomposition**:

$$\mathbf{L}\mathbf{L}^\top = \mathbf{K}_\theta + \sigma^2\mathbf{I}, \quad \text{where } \mathbf{L} \text{ is lower triangular.} \quad (6)$$

$$\log|\mathbf{K}_\theta + \sigma^2\mathbf{I}| = \log|\mathbf{L}\mathbf{L}^\top| = 2 \sum_{n=1}^N \log[\mathbf{L}]_{nn} \quad (7)$$

$$\mathbf{y}^\top (\mathbf{K}_\theta + \sigma^2\mathbf{I})^{-1} \mathbf{y} = \mathbf{y}^\top \mathbf{L}^\top{}^{-1} \mathbf{L}^{-1} \mathbf{y} \quad (8)$$

Cholesky decomposition is used for **practical implementation**.
(For the coursework, all implementations are fine.)

Computational complexity

1. Computing kernel matrix $O(N^2)$ time and $O(N^2)$ space
2. Eigendecomp and Cholesky are both often quoted to be $O(N^3)$ time.²
3. Logdet or inverse given the decomposition are fast
 - ▶ Logdet for both are $O(N)$
 - ▶ Inverse is $O(N^2)$

²Algorithms do exist with better asymptotic complexity, but they are definitely slower than $O(N^2)$.

Computational complexity

1. Computing kernel matrix $O(N^2)$ time and $O(N^2)$ space
2. Eigendecomposition and Cholesky are both often quoted to be $O(N^3)$ time.²
3. Logdet or inverse given the decomposition are fast
 - ▶ Logdet for both are $O(N)$
 - ▶ Inverse is $O(N^2)$

Can we take advantage of structure in the kernel matrix to do better?

²Algorithms do exist with better asymptotic complexity, but they are definitely slower than $O(N^2)$.

Low-rank kernels

For kernel matrices that are low-rank (have some zero eigenvalues),
i.e. with $M \ll N$

$$\mathbf{K} = \mathbf{P}\mathbf{P}^\top, \quad \text{where } \mathbf{K} \in \mathbb{R}^{N \times N}, \mathbf{P} \in \mathbb{R}^{N \times M} \quad (9)$$

³Can be proved by applying the Woodbury Matrix Identity, and the similar Matrix Determinant Lemma.

Low-rank kernels

For kernel matrices that are low-rank (have some zero eigenvalues),
i.e. with $M \ll N$

$$\mathbf{K} = \mathbf{P}\mathbf{P}^\top, \quad \text{where } \mathbf{K} \in \mathbb{R}^{N \times N}, \mathbf{P} \in \mathbb{R}^{N \times M} \quad (9)$$

we can compute the marginal likelihood more cheaply:³

$$\log|\mathbf{K} + \sigma^2\mathbf{I}_N| = \log|\sigma^2\mathbf{I}_N| + \log|\mathbf{I}_M + \sigma^{-2}\mathbf{P}^\top\mathbf{P}| \quad (10)$$

$$\mathbf{y}^\top(\mathbf{K} + \sigma^2\mathbf{I}_N)^{-1}\mathbf{y} = \mathbf{y}^\top\left(\sigma^{-2}\mathbf{I}_N - \sigma^{-2}\mathbf{P}(\sigma^2\mathbf{I}_M + \mathbf{P}^\top\mathbf{P})^{-1}\mathbf{P}^\top\right)\mathbf{y} \quad (11)$$

³Can be proved by applying the Woodbury Matrix Identity, and the similar Matrix Determinant Lemma.

Low-rank kernels

For kernel matrices that are low-rank (have some zero eigenvalues),
i.e. with $M \ll N$

$$\mathbf{K} = \mathbf{P}\mathbf{P}^\top, \quad \text{where } \mathbf{K} \in \mathbb{R}^{N \times N}, \mathbf{P} \in \mathbb{R}^{N \times M} \quad (9)$$

we can compute the marginal likelihood more cheaply:³

$$\log|\mathbf{K} + \sigma^2\mathbf{I}_N| = \log|\sigma^2\mathbf{I}_N| + \log|\mathbf{I}_M + \sigma^{-2}\mathbf{P}^\top\mathbf{P}| \quad (10)$$

$$\mathbf{y}^\top(\mathbf{K} + \sigma^2\mathbf{I}_N)^{-1}\mathbf{y} = \mathbf{y}^\top\left(\sigma^{-2}\mathbf{I}_N - \sigma^{-2}\mathbf{P}(\sigma^2\mathbf{I}_M + \mathbf{P}^\top\mathbf{P})^{-1}\mathbf{P}^\top\right)\mathbf{y} \quad (11)$$

Can be computed in $O(NM^2 + M^3)$.

³Can be proved by applying the Woodbury Matrix Identity, and the similar Matrix Determinant Lemma.

Decaying eigenvalues

For example, the squared exponential kernel

$$k(x, x') = \sigma_f^2 \exp\left(\frac{|x - x'|^2}{2\ell^2}\right), \quad \theta = \{\sigma_f, \ell\} \quad (12)$$

with 1D inputs i.i.d. from⁴ $\mathcal{N}(0, \sigma_D^2)$, we have that the eigenvalues of \mathbf{K} decay exponentially

$$\lambda_m = \sqrt{\frac{2a}{A}} B^m \quad \text{with } B < 1 \quad (13)$$

$$a^{-1} = 4\sigma_D^2, \quad b^{-1} = 2\ell^2, \quad c = \sqrt{a^2 + 2ab}, \quad A = a + b + c, \quad B = b/A$$

as the number of data $N \rightarrow \infty$.

⁴Conditions can be much weaker, but this is easier.

Low-rank approximations

So we know that for large N our \mathbf{K} will be **approximately low-rank**.
Can we approximate it with a low-rank matrix?

$$\mathbf{K} = \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^\top = [\mathbf{Q}_1 \quad \mathbf{Q}_2] \begin{bmatrix} \mathbf{\Lambda}_1 & \\ & \mathbf{\Lambda}_2 \end{bmatrix} \begin{bmatrix} \mathbf{Q}_1^\top \\ \mathbf{Q}_2^\top \end{bmatrix} = \mathbf{Q}_1\mathbf{\Lambda}_1\mathbf{Q}_1^\top + \mathbf{Q}_2\mathbf{\Lambda}_2\mathbf{Q}_2^\top \quad (14)$$

Low-rank approximations

So we know that for large N our \mathbf{K} will be **approximately low-rank**.
Can we approximate it with a low-rank matrix?

$$\mathbf{K} = \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^\top = [\mathbf{Q}_1 \quad \mathbf{Q}_2] \begin{bmatrix} \mathbf{\Lambda}_1 & \\ & \mathbf{\Lambda}_2 \end{bmatrix} \begin{bmatrix} \mathbf{Q}_1^\top \\ \mathbf{Q}_2^\top \end{bmatrix} = \mathbf{Q}_1\mathbf{\Lambda}_1\mathbf{Q}_1^\top + \mathbf{Q}_2\mathbf{\Lambda}_2\mathbf{Q}_2^\top \quad (14)$$

$$\hat{\mathbf{K}} = \mathbf{Q}_1\mathbf{\Lambda}_1\mathbf{Q}_1^\top + \cancel{\mathbf{Q}_2\mathbf{\Lambda}_2\mathbf{Q}_2^\top} \stackrel{\approx 0}{=} \mathbf{P}\mathbf{P}^\top, \quad \text{with } \mathbf{P} = \mathbf{Q}_1\mathbf{\Lambda}_1^{\frac{1}{2}}. \quad (15)$$

Low-rank approximations

So we know that for large N our \mathbf{K} will be **approximately low-rank**.
Can we approximate it with a low-rank matrix?

$$\mathbf{K} = \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^\top = [\mathbf{Q}_1 \quad \mathbf{Q}_2] \begin{bmatrix} \mathbf{\Lambda}_1 & \\ & \mathbf{\Lambda}_2 \end{bmatrix} \begin{bmatrix} \mathbf{Q}_1^\top \\ \mathbf{Q}_2^\top \end{bmatrix} = \mathbf{Q}_1\mathbf{\Lambda}_1\mathbf{Q}_1^\top + \mathbf{Q}_2\mathbf{\Lambda}_2\mathbf{Q}_2^\top \quad (14)$$

$$\hat{\mathbf{K}} = \mathbf{Q}_1\mathbf{\Lambda}_1\mathbf{Q}_1^\top + \cancel{\mathbf{Q}_2\mathbf{\Lambda}_2\mathbf{Q}_2^\top} \stackrel{\approx 0}{=} \mathbf{P}\mathbf{P}^\top, \quad \text{with } \mathbf{P} = \mathbf{Q}_1\mathbf{\Lambda}_1^{\frac{1}{2}}. \quad (15)$$

This gives the approximation with the minimal Frobenius norm:

$$\|\mathbf{K} - \hat{\mathbf{K}}\|_F = \sqrt{\sum_{i=M+1}^N \lambda_i} \quad (16)$$

Low-rank approximation: Data fit

Starting with the data-fit term:⁵

$$\begin{aligned}\mathbf{y}^\top (\mathbf{K} + \sigma^2 \mathbf{I})^{-1} \mathbf{y} &= \mathbf{y}^\top (\hat{\mathbf{K}} + \sigma^2 \mathbf{I} + \mathbf{K} - \hat{\mathbf{K}})^{-1} \mathbf{y} \\ &= \mathbf{y}^\top (\hat{\mathbf{K}} + \sigma^2 \mathbf{I})^{-1} \mathbf{y} - \underbrace{\mathbf{y}^\top (\hat{\mathbf{K}} + \sigma^2 \mathbf{I})^{-1} (\mathbf{K} - \hat{\mathbf{K}}) (\mathbf{K} + \sigma^2 \mathbf{I})^{-1} \mathbf{y}}_{\rightarrow 0 \text{ as } \|\mathbf{K} - \hat{\mathbf{K}}\|_F \rightarrow 0} \quad (17)\end{aligned}$$

⁵Applying the identity $(A + B)^{-1} = A^{-1} - A^{-1}B(A + B)^{-1}$.

Low-rank approximation: Data fit

Starting with the data-fit term:⁵

$$\begin{aligned}\mathbf{y}^\top (\mathbf{K} + \sigma^2 \mathbf{I})^{-1} \mathbf{y} &= \mathbf{y}^\top (\hat{\mathbf{K}} + \sigma^2 \mathbf{I} + \mathbf{K} - \hat{\mathbf{K}})^{-1} \mathbf{y} \\ &= \mathbf{y}^\top (\hat{\mathbf{K}} + \sigma^2 \mathbf{I})^{-1} \mathbf{y} - \underbrace{\mathbf{y}^\top (\hat{\mathbf{K}} + \sigma^2 \mathbf{I})^{-1} (\mathbf{K} - \hat{\mathbf{K}}) (\mathbf{K} + \sigma^2 \mathbf{I})^{-1} \mathbf{y}}_{\rightarrow 0 \text{ as } \|\mathbf{K} - \hat{\mathbf{K}}\|_F \rightarrow 0} \quad (17)\end{aligned}$$

You can place bounds on the error using knowledge of λ_M (the final included eigenvalue).

⁵Applying the identity $(A + B)^{-1} = A^{-1} - A^{-1}B(A + B)^{-1}$.

Low-rank approximation: Complexity penalty

And finishing with the complexity penalty:

$$\begin{aligned}\log|\mathbf{K} + \sigma^2\mathbf{I}| &= \sum_{n=1}^N \log(\lambda_n(\mathbf{K}) + \sigma^2) \\ &= \sum_{m=1}^M \log(\lambda_m(\mathbf{K}) + \sigma^2) + \sum_{r=M+1}^N \underbrace{\log(\lambda_r(\mathbf{K}) + \sigma^2)}_{\approx \sigma^2} \quad (18)\end{aligned}$$

$$\approx \sum_{m=1}^M \log(\lambda_m(\mathbf{K}) + \sigma^2) + \sum_{r=M+1}^N \log(\sigma^2) \quad (19)$$

$$= \log|\hat{\mathbf{K}} + \sigma^2\mathbf{I}| \quad (20)$$

Low-rank approximations overview

- ▶ Learning with exact matrix decompositions is expensive — $O(N^3)$ time, $O(N^2)$ memory
- ▶ Low-rank kernels improve things when $N \gg M$ — $O(NM^2)$ time, $O(NM)$ memory

Low-rank approximations overview

- ▶ Learning with exact matrix decompositions is expensive — $O(N^3)$ time, $O(N^2)$ memory
- ▶ Low-rank kernels improve things when $N \gg M$ — $O(NM^2)$ time, $O(NM)$ memory
- ▶ Low-rank approximations do exist!

Low-rank approximations overview

- ▶ Learning with exact matrix decompositions is expensive — $O(N^3)$ time, $O(N^2)$ memory
- ▶ Low-rank kernels improve things when $N \gg M$ — $O(NM^2)$ time, $O(NM)$ memory
- ▶ Low-rank approximations do exist!
- ▶ We used an impractical method (eigendecomposition) to find **which** low-rank columns to use, with a cost of $O(N^2M)$ time.

Low-rank approximations overview

- ▶ Learning with exact matrix decompositions is expensive — $O(N^3)$ time, $O(N^2)$ memory
- ▶ Low-rank kernels improve things when $N \gg M$ — $O(NM^2)$ time, $O(NM)$ memory
- ▶ Low-rank approximations do exist!
- ▶ We used an impractical method (eigendecomposition) to find **which** low-rank columns to use, with a cost of $O(N^2M)$ time.

For certain approximations, we have 1) fast methods for finding \mathbf{P} , and 2) proofs for how large M needs to be.

E.g. Burt et al [2019] show that for regression and the Squared Exponential kernel, we can get arbitrarily good approximations in

$$O(N(\log N)^{2D}(\log \log N)^2) = O(N^{1+\epsilon}), \quad \forall \epsilon > 0 \quad (21)$$

Practical constraints and implementation

GPs are slow, but often not even because of asymptotic complexity!

Practical constraints and implementation

GPs are slow, but often not even because of asymptotic complexity!

- ▶ Memory is often the main bottleneck (storing large matrices).

Practical constraints and implementation

GPs are slow, but often not even because of asymptotic complexity!

- ▶ Memory is often the main bottleneck (storing large matrices).
This is poorly suited to modern-day hardware

Practical constraints and implementation

GPs are slow, but often not even because of asymptotic complexity!

- ▶ Memory is often the main bottleneck (storing large matrices).
This is poorly suited to modern-day hardware (i.e. GPUs)

Practical constraints and implementation

GPs are slow, but often not even because of asymptotic complexity!

- ▶ Memory is often the main bottleneck (storing large matrices).
This is poorly suited to modern-day hardware (i.e. GPUs)
- ▶ Decompositions require serial computations.

Practical constraints and implementation

GPs are slow, but often not even because of asymptotic complexity!

- ▶ Memory is often the main bottleneck (storing large matrices).
This is poorly suited to modern-day hardware (i.e. GPUs)
- ▶ Decompositions require serial computations.
This is poorly suited to modern-day hardware

Practical constraints and implementation

GPs are slow, but often not even because of asymptotic complexity!

- ▶ Memory is often the main bottleneck (storing large matrices).
This is poorly suited to modern-day hardware (i.e. GPUs)
- ▶ Decompositions require serial computations.
This is poorly suited to modern-day hardware (i.e. GPUs)

Practical constraints and implementation

GPs are slow, but often not even because of asymptotic complexity!

- ▶ Memory is often the main bottleneck (storing large matrices).
This is poorly suited to modern-day hardware (i.e. GPUs)
- ▶ Decompositions require serial computations.
This is poorly suited to modern-day hardware (i.e. GPUs)
- ▶ Decompositions need high floating point precision.

Practical constraints and implementation

GPs are slow, but often not even because of asymptotic complexity!

- ▶ Memory is often the main bottleneck (storing large matrices).
This is poorly suited to modern-day hardware (i.e. GPUs)
- ▶ Decompositions require serial computations.
This is poorly suited to modern-day hardware (i.e. GPUs)
- ▶ Decompositions need high floating point precision.
This is poorly suited to modern-day hardware

Practical constraints and implementation

GPs are slow, but often not even because of asymptotic complexity!

- ▶ Memory is often the main bottleneck (storing large matrices).
This is poorly suited to modern-day hardware (i.e. GPUs)
- ▶ Decompositions require serial computations.
This is poorly suited to modern-day hardware (i.e. GPUs)
- ▶ Decompositions need high floating point precision.
This is poorly suited to modern-day hardware (i.e. GPUs)

There is hope

Approximations can make the computations look more similar to a neural network by

- ▶ Reduce the size of the matrix decomposition [Titsias, 2009]
- ▶ Improve data parallelism [Hensman et al, 2013]
- ▶ Increase reliance on fast matrix-matrix multiplications [Gardner et al, 2018; Wang et al, 2019]
- ▶ Remove matrix inverses altogether [van der Wilk et al, 2019]

There is hope

Approximations can make the computations look more similar to a neural network by

- ▶ Reduce the size of the matrix decomposition [Titsias, 2009]
- ▶ Improve data parallelism [Hensman et al, 2013]
- ▶ Increase reliance on fast matrix-matrix multiplications [Gardner et al, 2018; Wang et al, 2019]
- ▶ Remove matrix inverses altogether [van der Wilk et al, 2019]

This slide for reference only, not examinable.

Implementation matters!

We have good asymptotic bounds on the computational complexity!

Constants cause slowness \implies implementation matters

Implementation matters!

We have good asymptotic bounds on the computational complexity!

Constants cause slowness \implies implementation matters

E.g. *Exact Gaussian Processes on a Million Data Points*, Wang et al [2019] tailors a very good approximation to GPU hardware, with impressive results.

Conclusions on Computational Issues

- ▶ Gaussian processes are computationally expensive to train
- ▶ Low-rank approximations to the kernel matrix \mathbf{K} do exist
- ▶ Low-rank approximations improve the asymptotic complexity
- ▶ Still need work to optimise algorithms and tailor to modern hardware

Modelling limitations

We know that the generalisation capability of a model depends strongly on the properties of its prior.

Do we really have good priors?

Modelling limitations

We know that the generalisation capability of a model depends strongly on the properties of its prior.

Do we really have good priors?

Some criticisms...

Stationary kernels

We have seen many examples of **stationary** kernels:

$$k(\mathbf{x}, \mathbf{x}') = k(\mathbf{x} - \mathbf{x}') \quad (22)$$

e.g. the squared exponential.

- ▶ Same generalisation characteristic is applied through the entire input space.
- ▶ See `stationary-kernel-failure.ipynb` in the `inference-plots` GitHub repo.

Local kernels in high dimensions

Many kernels depend on some form of the Euclidean distance:

$$k(\mathbf{x}, \mathbf{x}') = \sigma_f^2 \exp\left[-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2\ell^2}\right] \quad (23)$$

This behaves badly in high dimensions (even MNIST, $D = 28^2 = 784$):

- ▶ Small differences in each dimension add up to large distances
→ low correlation
- ▶ Function is allowed to vary independently along each input dimension
- ▶ Prior with low lengthscale is too flexible, large lengthscale is too inflexible
- ▶ No middle ground

Conclusions

Gaussian processes have limitations for both

- ▶ low-dimensional inputs (e.g. stationarity)
- ▶ high-dimensional inputs (e.g. locality)

We need to either make much better priors, or incorporate GPs into more complex models.

References I

- [1] J. Gardner, G. Pleiss, K. Q. Weinberger, D. Bindel, and A. G. Wilson. Gpytorch: Blackbox matrix-matrix gaussian process inference with gpu acceleration. In S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett, editors, *Advances in Neural Information Processing Systems 31*, pages 7576–7586. Curran Associates, Inc., 2018.