Decision Theory

Mark van der Wilk

Department of Computing
Imperial College London

@markvanderwilk
m.vdwilk@imperial.ac.uk

February 13, 2023

Taking Action

What do you want to **use** a posterior for?

Taking Action

What do you want to **use** a posterior for?

- Posterior is a belief over what will happen.

Taking Action

What do you want to **use** a posterior for?

- Posterior is a belief over what will happen.
- You use beliefs to take action / make decisions.

Taking Action

What do you want to **use** a posterior for?

- Posterior is a belief over what will happen.
- You use beliefs to take action / make decisions.
- Uncertainty helps you to balance risk, reward, and new information.

Taking Action

What do you want to **use** a posterior for?

- Posterior is a belief over what will happen.
- You use beliefs to take action / make decisions.
- Uncertainty helps you to balance risk, reward, and new information.

We **learn** about the world,
so we can **act** in it,
to get outcomes that we **desire**.

Axiomatisation

As with uncertain reasoning, we can get maths to tell us how to make rational decisions, if we write down our wishlist of axioms.

Axiomatisation

As with uncertain reasoning, we can get maths to tell us how to make rational decisions, if we write down our wishlist of axioms.
Paraphrased axioms:

Axiomatisation

As with uncertain reasoning, we can get maths to tell us how to make rational decisions, if we write down our wishlist of axioms.
Paraphrased axioms:

1. Comparing probability distributions A and B over outcomes, you either prefer one over the other, or you have no preference.

Axiomatisation

As with uncertain reasoning, we can get maths to tell us how to make rational decisions, if we write down our wishlist of axioms.

Paraphrased axioms:

1. Comparing probability distributions A and B over outcomes, you either prefer one over the other, or you have no preference.
2. If you prefer A over B, and B over C, then you prefer A over C.

Axiomatisation

As with uncertain reasoning, we can get maths to tell us how to make rational decisions, if we write down our wishlist of axioms.
Paraphrased axioms:

1. Comparing probability distributions A and B over outcomes, you either prefer one over the other, or you have no preference.
2. If you prefer A over B, and B over C, then you prefer A over C.
3. Reduction of compound lotteries: If you prefer A over B, then you also prefer a 50% chance of getting A over a 50% chance of getting B, with a 50% chance of getting C.

Axiomatisation

As with uncertain reasoning, we can get maths to tell us how to make rational decisions, if we write down our wishlist of axioms.

Paraphrased axioms:

1. Comparing probability distributions A and B over outcomes, you either prefer one over the other, or you have no preference.
2. If you prefer A over B, and B over C, then you prefer A over C.
3. Reduction of compound lotteries: If you prefer A over B, then you also prefer a 50% chance of getting A over a 50% chance of getting B, with a 50% chance of getting C.

See Ch 16 “Artificial Intelligence: A Modern Approach”, Russell [1]

Von Neumann-Morgenstern Utility Theorem

- If you behave according to these axioms, then there exists a **utility function** that you maximise.

Von Neumann-Morgenstern Utility Theorem

- If you behave according to these axioms, then there exists a **utility function** that you maximise.
- In other words, if you believe you want to act according to those axioms, you should follow the principle of **Maximum Expected Utility**.

Von Neumann-Morgenstern Utility Theorem

- If you behave according to these axioms, then there exists a **utility function** that you maximise.
- In other words, if you believe you want to act according to those axioms, you should follow the principle of **Maximum Expected Utility**.

Principle of Maximum Expected Utility

1. Define a utility function $U : \mathcal{X} \times \mathcal{A} \to \mathbb{R}$
   $\mathcal{X}$: Space of outcomes. $\mathcal{A}$: Space of actions. Quantifies how good an outcome is, if you take a particular action.

Principle of Maximum Expected Utility

1. Define a utility function $U : \mathcal{X} \times \mathcal{A} \to \mathbb{R}$
   $\mathcal{X}$: Space of outcomes. $\mathcal{A}$: Space of actions. Quantifies how good an outcome is, if you take a particular action.
2. Compute expected utility for your actions.
   Your beliefs are a distribution over outcomes given action.

$$u(a) = \mathbb{E}_{p(x|\mathcal{D},a)}[U(x,a)] \tag{1}$$

Principle of Maximum Expected Utility

1. Define a utility function $U : \mathcal{X} \times \mathcal{A} \to \mathbb{R}$
   $\mathcal{X}$: Space of outcomes. $\mathcal{A}$: Space of actions. Quantifies how good an outcome is, if you take a particular action.
2. Compute expected utility for your actions.
   Your beliefs are a distribution over outcomes given action.

$$u(a) = \mathbb{E}_{p(x|\mathcal{D},a)}[U(x,a)] \tag{1}$$

3. At the time of decision making, choose action which maximises expected utility.

$$a^* = \operatorname*{argmax}_a u(a) \tag{2}$$

Exercise: Why do we predict the mean?

Your regression model gives you $p(y|X, \mathbf{y})$. You need to give a "best guess" y_p, and your utility will be $U(y, y_p) = -(y - y_p)^2$.

▶▶ Find y_p in terms of properties of $p(y|X, \mathbf{y})$.

Exercise: Why do we predict the mean?

Your regression model gives you $p(y|X, \mathbf{y})$. You need to give a "best guess" y_p, and your utility will be $U(y, y_p) = -(y - y_p)^2$.

▶▶ Find y_p in terms of properties of $p(y|X, \mathbf{y})$.

Solution: We want $\frac{\mathrm{d}}{\mathrm{d}y_p}u(y_p) = \frac{\mathrm{d}}{\mathrm{d}y_p}\mathbb{E}_{p(y|X,\mathbf{y})}\left[U(y, y_p)\right] = 0.$

Exercise: Why do we predict the mean?

Your regression model gives you $p(y|X, \mathbf{y})$. You need to give a "best guess" y_p, and your utility will be $U(y, y_p) = -(y - y_p)^2$.

▶▶ Find y_p in terms of properties of $p(y|X, \mathbf{y})$.

Solution: We want $\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) = \frac{\mathrm{d}}{\mathrm{d}y_p} \mathbb{E}_{p(y|X,\mathbf{y})}\left[U(y, y_p)\right] = 0$.

$$\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) = \frac{\mathrm{d}}{\mathrm{d}y_p} \int_{-\infty}^{\infty} p(y|X, \mathbf{y}) U(y, y_p) \mathrm{d}y$$

Exercise: Why do we predict the mean?

Your regression model gives you $p(y|X, \mathbf{y})$. You need to give a "best guess" y_p, and your utility will be $U(y, y_p) = -(y - y_p)^2$.

Find y_p in terms of properties of $p(y|X, \mathbf{y})$.

Solution: We want $\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) = \frac{\mathrm{d}}{\mathrm{d}y_p} \mathbb{E}_{p(y|X,\mathbf{y})}\left[U(y, y_p)\right] = 0$.

$$\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) = \frac{\mathrm{d}}{\mathrm{d}y_p} \int_{-\infty}^{\infty} p(y|X, \mathbf{y}) U(y, y_p) \mathrm{d}y$$

$$= -\int p(y|X, \mathbf{y}) \frac{\partial}{\partial y_p} (y - y_p)^2 \mathrm{d}y \qquad \text{Leibniz Integral Rule}$$

Exercise: Why do we predict the mean?

Your regression model gives you $p(y|X, \mathbf{y})$. You need to give a "best guess" y_p, and your utility will be $U(y, y_p) = -(y - y_p)^2$.

⏩ Find y_p in terms of properties of $p(y|X, \mathbf{y})$.

Solution: We want $\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) = \frac{\mathrm{d}}{\mathrm{d}y_p} \mathbb{E}_{p(y|X,\mathbf{y})}\left[U(y, y_p)\right] = 0.$

$$
\begin{aligned}
\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) &= \frac{\mathrm{d}}{\mathrm{d}y_p} \int_{-\infty}^{\infty} p(y|X, \mathbf{y}) U(y, y_p) \mathrm{d}y \\
&= -\int p(y|X, \mathbf{y}) \frac{\partial}{\partial y_p} (y - y_p)^2 \mathrm{d}y && \text{Leibniz Integral Rule} \\
&= 2 \int p(y|X, \mathbf{y}) (y - y_p) \mathrm{d}y
\end{aligned}
$$

Exercise: Why do we predict the mean?

Your regression model gives you $p(y|X, \mathbf{y})$. You need to give a "best guess" y_p, and your utility will be $U(y, y_p) = -(y - y_p)^2$.

▶▶ Find y_p in terms of properties of $p(y|X, \mathbf{y})$.

Solution: We want $\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) = \frac{\mathrm{d}}{\mathrm{d}y_p} \mathbb{E}_{p(y|X,\mathbf{y})}\left[U(y, y_p)\right] = 0.$

$$
\begin{aligned}
\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) &= \frac{\mathrm{d}}{\mathrm{d}y_p} \int_{-\infty}^{\infty} p(y|X, \mathbf{y}) U(y, y_p) \mathrm{d}y \\
&= -\int p(y|X, \mathbf{y}) \frac{\partial}{\partial y_p} (y - y_p)^2 \mathrm{d}y && \text{Leibniz Integral Rule} \\
&= 2 \int p(y|X, \mathbf{y})(y - y_p) \mathrm{d}y \\
0 &= 2\mathbb{E}_{p(y|X,\mathbf{y})}[y] - 2y_p
\end{aligned}
$$

Exercise: Why do we predict the mean?

Your regression model gives you $p(y|X, \mathbf{y})$. You need to give a "best guess" y_p, and your utility will be $U(y, y_p) = -(y - y_p)^2$.

▶▶ Find y_p in terms of properties of $p(y|X, \mathbf{y})$.

Solution: We want $\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) = \frac{\mathrm{d}}{\mathrm{d}y_p} \mathbb{E}_{p(y|X,\mathbf{y})}[U(y, y_p)] = 0$.

$$
\begin{aligned}
\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) &= \frac{\mathrm{d}}{\mathrm{d}y_p} \int_{-\infty}^{\infty} p(y|X, \mathbf{y}) U(y, y_p) \mathrm{d}y \\
&= -\int p(y|X, \mathbf{y}) \frac{\partial}{\partial y_p} (y - y_p)^2 \mathrm{d}y && \text{Leibniz Integral Rule} \\
&= 2 \int p(y|X, \mathbf{y})(y - y_p) \mathrm{d}y \\
0 &= 2\mathbb{E}_{p(y|X,\mathbf{y})}[y] - 2y_p \\
y_p &= \mathbb{E}_{p(y|X,\mathbf{y})}[y], \qquad \text{i.e. } y_p \text{ should be the } \textbf{mean}!
\end{aligned}
$$

Exercise: Absolute error loss.

Your regression model gives you $p(y|X, \mathbf{y})$. You need to give a "best guess" y_p, and your utility will be $U(y, y_p) = -|y - y_p|$.

▶▶ Find y_p in terms of properties of $p(y|X, \mathbf{y})$.

Exercise: Absolute error loss.

Your regression model gives you $p(y|X, \mathbf{y})$. You need to give a "best guess" y_p, and your utility will be $U(y, y_p) = -|y - y_p|$.

▶▶ Find y_p in terms of properties of $p(y|X, \mathbf{y})$.

Solution: We want $\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) = \frac{\mathrm{d}}{\mathrm{d}y_p} \mathbb{E}_{p(y|X,\mathbf{y})}\left[U(y, y_p)\right] = 0$.

Exercise: Absolute error loss.

Your regression model gives you $p(y|X, \mathbf{y})$. You need to give a "best guess" y_p, and your utility will be $U(y, y_p) = -|y - y_p|$.

⏩ Find y_p in terms of properties of $p(y|X, \mathbf{y})$.

Solution: We want $\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) = \frac{\mathrm{d}}{\mathrm{d}y_p} \mathbb{E}_{p(y|X,\mathbf{y})}\left[U(y, y_p)\right] = 0$.

$$\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) = \frac{\mathrm{d}}{\mathrm{d}y_p} \int_{-\infty}^{\infty} p(y|X, \mathbf{y}) U(y, y_p) \mathrm{d}y$$

Exercise: Absolute error loss.

Your regression model gives you $p(y|X, \mathbf{y})$. You need to give a "best guess" y_p, and your utility will be $U(y, y_p) = -|y - y_p|$.

▶▶ Find y_p in terms of properties of $p(y|X, \mathbf{y})$.

Solution: We want $\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) = \frac{\mathrm{d}}{\mathrm{d}y_p} \mathbb{E}_{p(y|X,\mathbf{y})}\left[U(y, y_p)\right] = 0$.

$$
\begin{aligned}
\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) &= \frac{\mathrm{d}}{\mathrm{d}y_p} \int_{-\infty}^{\infty} p(y|X, \mathbf{y}) U(y, y_p) \mathrm{d}y \\
&= \int p(y|X, \mathbf{y}) \frac{\partial}{\partial y_p} U(y, y_p) \mathrm{d}y && \text{Leibniz Integral Rule}
\end{aligned}
$$

Exercise: Absolute error loss.

Your regression model gives you $p(y|X, \mathbf{y})$. You need to give a "best guess" y_p, and your utility will be $U(y, y_p) = -|y - y_p|$.

Find y_p in terms of properties of $p(y|X, \mathbf{y})$.

Solution: We want $\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) = \frac{\mathrm{d}}{\mathrm{d}y_p} \mathbb{E}_{p(y|X,\mathbf{y})}\left[U(y, y_p)\right] = 0.$

$$\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) = \frac{\mathrm{d}}{\mathrm{d}y_p} \int_{-\infty}^{\infty} p(y|X, \mathbf{y}) U(y, y_p) \mathrm{d}y$$

$$= \int p(y|X, \mathbf{y}) \frac{\partial}{\partial y_p} U(y, y_p) \mathrm{d}y \qquad \text{Leibniz Integral Rule}$$

$$= \int p(y|X, \mathbf{y}) \begin{Bmatrix} -1 & \text{if } y > y_p \\ 1 & \text{if } y < y_p \end{Bmatrix} \mathrm{d}y$$

Exercise: Absolute error loss.

Your regression model gives you $p(y|X, \mathbf{y})$. You need to give a "best guess" y_p, and your utility will be $U(y, y_p) = -|y - y_p|$.

⏩ Find y_p in terms of properties of $p(y|X, \mathbf{y})$.

Solution: We want $\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) = \frac{\mathrm{d}}{\mathrm{d}y_p} \mathbb{E}_{p(y|X,\mathbf{y})}\left[U(y, y_p)\right] = 0$.

$$
\begin{aligned}
\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) &= \frac{\mathrm{d}}{\mathrm{d}y_p} \int_{-\infty}^{\infty} p(y|X, \mathbf{y}) U(y, y_p) \mathrm{d}y \\
&= \int p(y|X, \mathbf{y}) \frac{\partial}{\partial y_p} U(y, y_p) \mathrm{d}y && \text{Leibniz Integral Rule} \\
&= \int p(y|X, \mathbf{y}) \begin{Bmatrix} -1 & \text{if } y > y_p \\ 1 & \text{if } y < y_p \end{Bmatrix} \mathrm{d}y \\
&= -\int_{-\infty}^{y_p} p(y|X, \mathbf{y}) \mathrm{d}y + \int_{y_p}^{\infty} p(y|X, \mathbf{y}) \mathrm{d}y
\end{aligned}
$$

Exercise: Absolute error loss.

Your regression model gives you $p(y|X, \mathbf{y})$. You need to give a "best guess" y_p, and your utility will be $U(y, y_p) = -|y - y_p|$.

⏩ Find y_p in terms of properties of $p(y|X, \mathbf{y})$.

Solution: We want $\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) = \frac{\mathrm{d}}{\mathrm{d}y_p} \mathbb{E}_{p(y|X,\mathbf{y})}\left[U(y, y_p)\right] = 0$.

$$
\begin{aligned}
\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) &= \frac{\mathrm{d}}{\mathrm{d}y_p} \int_{-\infty}^{\infty} p(y|X, \mathbf{y}) U(y, y_p) \mathrm{d}y \\
&= \int p(y|X, \mathbf{y}) \frac{\partial}{\partial y_p} U(y, y_p) \mathrm{d}y && \text{Leibniz Integral Rule} \\
&= \int p(y|X, \mathbf{y}) \begin{Bmatrix} -1 & \text{if } y > y_p \\ 1 & \text{if } y < y_p \end{Bmatrix} \mathrm{d}y \\
&= -\int_{-\infty}^{y_p} p(y|X, \mathbf{y}) \mathrm{d}y + \int_{y_p}^{\infty} p(y|X, \mathbf{y}) \mathrm{d}y \\
0 &= -F(y_p) + (1 - F(y_p)) && F \text{ is cdf}
\end{aligned}
$$

Exercise: Absolute error loss.

Your regression model gives you $p(y|X, \mathbf{y})$. You need to give a "best guess" y_p, and your utility will be $U(y, y_p) = -|y - y_p|$.

⏩ Find y_p in terms of properties of $p(y|X, \mathbf{y})$.

Solution: We want $\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) = \frac{\mathrm{d}}{\mathrm{d}y_p} \mathbb{E}_{p(y|X,\mathbf{y})}\left[U(y, y_p)\right] = 0.$

$$
\begin{aligned}
\frac{\mathrm{d}}{\mathrm{d}y_p} u(y_p) &= \frac{\mathrm{d}}{\mathrm{d}y_p} \int_{-\infty}^{\infty} p(y|X, \mathbf{y}) U(y, y_p) \mathrm{d}y \\
&= \int p(y|X, \mathbf{y}) \frac{\partial}{\partial y_p} U(y, y_p) \mathrm{d}y && \text{Leibniz Integral Rule} \\
&= \int p(y|X, \mathbf{y}) \begin{Bmatrix} -1 & \text{if } y > y_p \\ 1 & \text{if } y < y_p \end{Bmatrix} \mathrm{d}y \\
&= -\int_{-\infty}^{y_p} p(y|X, \mathbf{y}) \mathrm{d}y + \int_{y_p}^{\infty} p(y|X, \mathbf{y}) \mathrm{d}y \\
0 &= -F(y_p) + (1 - F(y_p)) && F \text{ is cdf}
\end{aligned}
$$

$\implies y_p = F^{-1}(0.5)\,,$ i.e. y_p should be the **median**!

Applications and Applicability

Hugely influential theory:

- Philosophy: Utilitarianism effective altruism
- Psychology: Rational behaviour as a model for human behaviour.
- Economics: How to optimise investments.
- Game theory: Analyse implications of rational choice.
- Politics: Voting systems (Arrow's impossibility theorems).

Applications and Applicability

Hugely influential theory:

- Philosophy: Utilitarianism effective altruism
- Psychology: Rational behaviour as a model for human behaviour.
- Economics: How to optimise investments.
- Game theory: Analyse implications of rational choice.
- Politics: Voting systems (Arrow's impossibility theorems).

Discussions about desirability/realism of axioms:

- Naïve application leads to bad behaviour (St Petersburg Paradox)
- Can you express desires as utility functions?
- The Law of Perverse Optimization: Whenever a desired behaviour is formulated as a utility, optimising for the utility will give you behaviour you didn't want. (AI Paperclip factory)
- Human's don't behave according to the axioms. But perhaps for good reason?
- Bounded rationality: It assumes you can compute the optimal decision.

Conclusion

- Decision theory is easy (in principle): See slide 5.

Conclusion

- Decision theory is easy (in principle): See slide 5.
- The implications of decision theory are vast, but sadly we don't have time for this.

Conclusion

- Decision theory is easy (in principle): See slide 5.
- The implications of decision theory are vast, but sadly we don't have time for this.
- Next: Things that are easy in principle, are often very hard in practice.

References I

[1] S. J. Russell. *Artificial intelligence a modern approach*. Pearson Education, Inc., 2010.